

## Data Modelling 4:

# Transactions Overview

### **Definition:**

A transaction in SQL is a sequence of one or more SQL statements executed as a single unit. The key concept is that all statements in a transaction are treated as a whole—either they all succeed (commit) or they all fail (rollback). This ensures data integrity and consistency.

---

## Transaction Access Modes

### 1. READ-ONLY:

- **Definition:** Restricts the transaction to perform only SELECT statements. No data modifications are allowed (INSERT, UPDATE, DELETE).
- **Purpose:** Ensures data is not accidentally modified, allowing unlimited reads.
- **Use Case:** Ideal for reporting, data analytics, or scenarios where no data alteration is needed.

#### **SQL:**

```
START TRANSACTION READ ONLY;
```

```
-- Fetch some data (this is allowed in READ ONLY mode)
```

```
SELECT * FROM employees;
```

```
INSERT INTO employees (employee_id, employee_name, department_id,  
manager_id, salary) VALUES (1, 'John Doe', 101, 10, 50000);
```

```
COMMIT;
```

### 2. READ-WRITE:

- **Definition:** Allows the transaction to perform both read (SELECT) and write (INSERT, UPDATE, DELETE) operations. This is the default mode.
- **Purpose:** Enables data modifications within the transaction.

- **Use Case:** Suitable for typical database operations where data changes are necessary.

**SQL:**

```
START TRANSACTION READ WRITE;
```

```
-- Fetch data from the table (allowed in READ-WRITE mode)
```

```
SELECT * FROM employees;
```

```
INSERT INTO employees (employee_id, employee_name, department_id, manager_id, salary) VALUES (12, 'John Doe', 101, 10, 50000);
```

```
UPDATE employees SET salary = 55000 WHERE employee_id = 1;
```

```
COMMIT;
```

---

## Transaction Isolation Levels

### 1. Read Uncommitted:

- **Scenario:** Transaction 1 modifies data, and Transaction 2 reads the uncommitted changes before Transaction 1 commits or rolls back.
- **Outcome:** Transaction 2 may see "dirty data" that could be rolled back, leading to inconsistencies.

**SQL:**

```
TRANSACTION 1:
```

```
SELECT * FROM accounts;
```

```
SET TRANSACTION ISOLATION LEVEL READ UNCOMMITTED;
```

```
START TRANSACTION;
```

```
UPDATE accounts SET balance = balance - 500 WHERE user_id = 1;
```

```
-- Start Transaction 2
```

```
SELECT * FROM accounts;
```

```
ROLLBACK;
```

#### **TRANSACTION 2:**

```
SELECT * FROM accounts;  
  
SET TRANSACTION ISOLATION LEVEL READ UNCOMMITTED;  
  
START TRANSACTION;  
  
SELECT * FROM accounts; //uncommitted changes are getting reflected here  
  
-- End Transaction 1  
  
COMMIT;  
  
SELECT * FROM accounts; //committed changes are getting reflected here as transaction1 is executed
```

## **2. Read Committed**

- **Scenario:** Transaction 1 modifies data but has not yet committed. Transaction 2 can only see data committed by other transactions.
- **Outcome:** Prevents dirty reads. Transaction 2 reads only committed data, ensuring data consistency.

**SQL:**

#### **TRANSACTION 1:**

```
SELECT * FROM accounts;  
  
SET TRANSACTION ISOLATION LEVEL READ COMMITTED;  
  
START TRANSACTION;  
  
UPDATE accounts SET balance = balance - 500 WHERE user_id = 1;  
  
-- Start Transaction 2  
  
SELECT * FROM accounts;  
  
COMMIT;
```

#### **TRANSACTION 2:**

```
SELECT * FROM accounts;  
  
SET TRANSACTION ISOLATION LEVEL READ UNCOMMITTED;
```

```
START TRANSACTION;
```

```
SELECT * FROM accounts; //only committed changes are getting reflected here
```

```
-- End Transaction 1
```

```
COMMIT;
```

```
SELECT * FROM accounts; //committed changes are getting reflected
```

### 3. Repeatable Read

- **Scenario:** Transaction 1 reads a record, then Transaction 2 modifies the same record. Transaction 1 reads the record again and sees the initial value, even if Transaction 2 committed changes.
- **Outcome:** Prevents non-repeatable reads. Transaction 1 will continue to see the old value, preserving read consistency.

```
SQL:
```

```
// Only Committed values are being read
```

```
TRANSACTION 1:
```

```
SET TRANSACTION ISOLATION LEVEL REPEATABLE READ;
```

```
START TRANSACTION;
```

```
UPDATE accounts SET balance = balance + 500 WHERE user_id = 1;
```

```
SELECT * FROM accounts;
```

```
-- Start Transaction 2
```

```
COMMIT;
```

```
TRANSACTION 2:
```

```
-- Reads the original value since Transaction 1 hasn't committed yet
```

```
SET TRANSACTION ISOLATION LEVEL REPEATABLE READ;
```

```
SELECT balance FROM accounts WHERE user_id = 1;
```

```
COMMIT;
```

// Once read, the snapshot is being maintained through out the transaction  
**TRANSACTION 1.1:**

**SET TRANSACTION ISOLATION LEVEL REPEATABLE READ;**

**START TRANSACTION;**

-- Start Transaction 2.1

**UPDATE accounts SET balance = balance + 500 WHERE user\_id = 1;**

**SELECT \* FROM accounts;**

**COMMIT;**

**TRANSACTION 2.1:**

-- Reads the original value from DB and retains it through out the Transaction

**SET TRANSACTION ISOLATION LEVEL REPEATABLE READ;**

**SELECT balance FROM accounts WHERE user\_id = 1; //read the original DB value**

-- Execute Transaction 1.1

**SELECT balance FROM accounts WHERE user\_id = 1; //still showing original value although transaction 1.1 has committed the new value in DB**

**COMMIT;**

// Blocking Phantom Reads in Repeatable Read mode

**TRANSACTION 1.2:**

**SET TRANSACTION ISOLATION LEVEL REPEATABLE READ;**

**START TRANSACTION;**

**SELECT \* FROM accounts WHERE department\_id = 101;**

-- do one more transaction which inserts some new data in this table with department\_id 101

-- Start Transaction 2.2

**SELECT \* FROM accounts WHERE department\_id = 101; // data is getting reflected here**

**COMMIT;**

## **TRANSACTION 2.2:**

-- New insertion is gonna be made in this transaction, but the data will not be reflected in another already running txn.

**SET TRANSACTION ISOLATION LEVEL REPEATABLE READ;**

**START TRANSACTION;**

**INSERT INTO accounts (user\_id, user\_name, account\_type, balance, department\_id)**

**VALUES (7, 'Sachin Tendulkar', 'All Rounder', 30000, 101);**

**COMMIT;**

**-- check the data in already running Transaction 1.2**

## **4. Serializable**

- **Scenario:** Transaction 1 queries a set of rows, then Transaction 2 inserts a new row matching the query. Transaction 1 does not see the new row, preventing phantom reads.
- **Outcome:** Provides the strictest isolation, ensuring no phantom reads occur by completely isolating the transaction.

**SQL:**

**// Only Committed values are being read**

**TRANSACTION 1:**

**SET TRANSACTION ISOLATION LEVEL SERIALIZABLE;**

**START TRANSACTION;**

**SELECT \* FROM accounts WHERE department\_id = 101;**

**-- This query will lock all rows with department\_id = 101**

**-- and also prevent any other inserts, updates, or deletes in this range.**

**-- Start Transaction 2**

**COMMIT;**

**TRANSACTION 2:**

```
SET TRANSACTION ISOLATION LEVEL SERIALIZABLE;  
  
START TRANSACTION;  
  
INSERT INTO accounts (user_id, user_name, account_type, balance,  
department_id)  
  
VALUES (8, 'VVS LAXMAN', 'Batting', 20000, 101);  
  
// its gonna wait as lock has been taken by transaction 1 which is not yet completed.  
  
ROLLBACK;  
  
SELECT * FROM accounts;
```

---

## Rollbacks and Cascading Rollbacks

- **Rollback:** Reverts all changes made during a transaction, preserving database consistency if an error occurs. Rollbacks ensure the database adheres to the ACID principles.
  - **Cascading Rollbacks:** Occur when the rollback of one transaction triggers rollbacks in other dependent transactions, maintaining overall consistency.
- 

## Autocommit Modes

### 1. Autocommit ON:

- **Definition:** Each SQL statement is a separate transaction that is immediately committed upon success.
- **Outcome:** If a statement fails, only the unsuccessful statement is rolled back, while prior successful statements remain committed.

**SQL:**

**SESSION 1:**

```
SET autocommit = 1;
```

```
INSERT INTO accounts (user_id, user_name, account_type, balance, department_id) VALUES (15, 'Rinku Singh', 'Batsman', 25000, 106);
```

```
UPDATE accounts SET balance = 75000 WHERE user_id = 1;
```

-- Start Another session to read from accounts table

```
select * from accounts;
```

**SESSION 2:**

```
select * from accounts; // reads updated values as auto commit is on
```

## 2. Autocommit OFF:

- **Definition:** Transactions require explicit management with `BEGIN TRANSACTION`, `COMMIT`, and `ROLLBACK` commands.
- **Outcome:** Ensures that if one statement fails, the entire transaction can be rolled back, preventing partial commits.

**SQL:**

**SESSION 1:**

```
SET autocommit = 0;
```

```
UPDATE accounts SET balance = 50000 WHERE user_id = 1;
```

-- Start Another session to read from accounts table

```
select * from accounts;
```

```
COMMIT;
```

**SESSION 2:**

```
select * from accounts; // reads older value as Session1's update is not committed yet
```

**NOTE:** After setting auto commit, if we start the transaction its gonna over ride the auto commit config, as transaction starts with START TRANSACTION and ends with either COMMIT or ROLLBACK.

---

## Transactional Logs in MySQL

MySQL uses three main types of logs to ensure data durability and consistency during transactions:

### 1. BinLog (Binary Log)

- Logs modifications from INSERT, UPDATE, and DELETE statements, but only for committed transactions.
- **Purpose:** Used for replication and point-in-time recovery, not for rollbacks.

### 2. Redo Log

- Stores changes in a redo log buffer before writing to disk.
- **Purpose:** Assists with crash recovery by applying committed changes not yet written to data files.

### 3. Undo Log

- Maintains "before" versions of data for rollback purposes.
  - **Purpose:** Supports Multi-Version Concurrency Control (MVCC), allowing consistent snapshots for transactions.
- 

## Locking Mechanisms:

### 1. Pessimistic Locking

- **Definition:** Locks data immediately to prevent other transactions from reading or modifying it.
- **Implementation:**
  - **SELECT ... FOR UPDATE:** Exclusive lock, blocking both reads and writes from other transactions.
  - **LOCK IN SHARE MODE/ SELECT ... FOR SHARE:** Shared lock, allowing reads but blocking writes.
- **Pros:** Guarantees consistency by blocking conflicting operations.
- **Cons:** Can lead to performance issues and deadlocks in high-concurrency environments.

**SQL:**

**SELECT... FOR UPDATE:**

**SESSION 1:**

**SELECT \* FROM accounts WHERE user\_id = 1;**

-- Set isolation level to REPEATABLE READ for the transaction

**SET TRANSACTION ISOLATION LEVEL REPEATABLE READ;**

**START TRANSACTION;**

-- Start Another session to do operation on this table.

**SELECT \* FROM accounts WHERE department\_id = 101 LIMIT 1 FOR UPDATE;**

**COMMIT;**

**SESSION 2: as we are only reading the lock will not be applied**

**SET TRANSACTION ISOLATION LEVEL REPEATABLE READ;**

**START TRANSACTION;**

**SELECT \* FROM accounts WHERE department\_id = 101;**

**COMMIT;**

**SESSION 2.1: lock will be applied as we are reding in Select.. For Update mode**

**SET TRANSACTION ISOLATION LEVEL REPEATABLE READ;**

**START TRANSACTION;**

**SELECT \* FROM accounts WHERE department\_id = 101 limit 1 FOR UPDATE;**

**COMMIT;**

**SESSION 2.2: lock will be applied as we are updating the data**

**SET TRANSACTION ISOLATION LEVEL REPEATABLE READ;**

**START TRANSACTION;**

```
UPDATE accounts SET balance = 50000 WHERE user_id = 1;  
COMMIT;
```

**SELECT... FOR SHARE:**

**SESSION 1:**

```
SET TRANSACTION ISOLATION LEVEL REPEATABLE READ;  
START TRANSACTION;  
SELECT * FROM accounts WHERE department_id = 101 LIMIT 1 FOR SHARE;  
-- Start Another session to do operation on this table.  
COMMIT;
```

**SESSION 2: as we are only reading the lock will not be applied**

```
SET TRANSACTION ISOLATION LEVEL REPEATABLE READ;  
START TRANSACTION;  
SELECT * FROM accounts WHERE department_id = 101;  
COMMIT;
```

**SESSION 2.1: lock will not be applied as we are reading in share mode.**

```
SET TRANSACTION ISOLATION LEVEL REPEATABLE READ;  
START TRANSACTION;  
SELECT * FROM accounts WHERE department_id = 101 limit 1 FOR SHARE;  
COMMIT;
```

**SESSION 2.2: lock will be applied as we are updating the data**

```
SET TRANSACTION ISOLATION LEVEL REPEATABLE READ;  
START TRANSACTION;
```

```
UPDATE accounts SET balance = 50000 WHERE department_id = 101;  
ROLLBACK;
```

## 2. Optimistic Locking:

- **Definition:** Assumes conflicts are rare, checking for conflicts only at commit time.
- **Process:**
  1. **Read:** The transaction reads data and its version/timestamp.
  2. **Update:** Checks for version changes before committing.
  3. **Commit:** Proceeds if no changes were detected; otherwise, retries.
- **Pros:** High concurrency and low blocking.
- **Cons:** Increased retries in conflict-heavy environments.

**NOTE:** SQL do not implement this feature, this feature gets implemented either in application layer or in query layer.

**SQL:**

**SESSION 1:**

```
START TRANSACTION;
```

```
SELECT * FROM accounts WHERE id = 1;
```

```
-- Update the balance if the version is still 0
```

```
UPDATE accounts SET balance = balance + 500, version = version + 1 WHERE id = 1  
AND version = 0;
```

```
SELECT * FROM accounts WHERE id = 1;
```

**-- Start Another session to do operation on this table.**

**-- Commit the transaction**

```
COMMIT;
```

**SESSION 2: as we are only reading the lock will not be applied**

```
START TRANSACTION;
```

```
SELECT * FROM accounts WHERE id = 1;
```

```
-- Commit Transaction 1
```

```
UPDATE accounts SET balance = balance + 10000, version = version + 1 WHERE id = 1 AND version = 1;
```

```
select * from accounts where id = 1;
```

```
COMMIT;
```

```
select * from accounts where id = 1;
```

```
// this transaction's update will not be reflected in DB as version 1 is not there at the time of executing this update in DB the version was 2.
```

---

## Gap Locks and Next Key Locks

- **Gap Locks:** Lock gaps between rows in index ranges, preventing insertions to avoid phantom reads.
    - Example: In a table with entries (1, 3, 5), a SELECT on values > 1 would lock gaps between these rows to block insertions at 2 and 4.
  - **Next Key Locks:** Lock both the current row and the gap before it, blocking changes within the scanned range.
    - Example: Selecting row 3 locks rows and gaps between 1 and 3, preventing changes that could affect future reads.
- 

## Auto Increment Lock Modes

1. **Mode 0 (Traditional):** Table-level lock for the full insert, blocking other insertions.
2. **Mode 1 (Increment Locking):** Short table lock only while generating the next ID.
3. **Mode 2 (No Locking):** Lightweight mutex with no table or row locks, allowing higher concurrency.

---

## Locking and Isolation Levels Mapping

Isolation Level	Locking Mechanism	Prevents
<b>READ UNCOMMITTED</b>	No locks	Nothing (allows dirty reads)
<b>READ COMMITTED</b>	Shared for reads, exclusive for writes	Dirty reads
<b>REPEATABLE READ</b>	Shared, exclusive, gap, and next-key locks	Dirty reads, non-repeatable, phantom reads
<b>SERIALIZABLE</b>	Full isolation with range locks	Dirty, non-repeatable, phantom reads, full serializability

### TABLES REQUIRED:

-- Creating accounts table

```
CREATE TABLE accounts (
    id INT PRIMARY KEY AUTO_INCREMENT,
    user_id INT UNIQUE KEY NOT NULL,
    user_name VARCHAR(50) NOT NULL,
    account_type VARCHAR(20) NOT NULL,
    balance INT(20),
    department_id VARCHAR(20) NOT NULL,
);
```

CREATE TABLE employees (

```
employee_id INT PRIMARY KEY,
employee_name VARCHAR(50) NOT NULL,
department_id INT,
manager_id INT,
salary INT
);
```

```
ALTER TABLE accounts ADD COLUMN version INT DEFAULT 0;
```