

Event-Driven Architecture & Messaging

Event-Driven Architecture (EDA) is a design pattern where services communicate by producing and consuming **events**. An event is a record of a significant change in state (e.g., `order_placed`, `user_signed_up`). This model promotes loose coupling, as services don't need to know about each other; they only care about the events.

- **Observer/Subscriber Pattern:** This is the core of EDA. **Producers** (or Publishers) generate events without knowing who will receive them. **Consumers** (or Subscribers) express interest in certain types of events and are notified when they occur.
- **Message Queues (Kafka, RabbitMQ, GCP Pub/Sub):** These are the backbone of EDA, acting as a durable broker between producers and consumers. They receive events from producers and hold them until consumers are ready to process them, enabling asynchronous communication.
- **Messaging Models:**
 - **Push vs. Pull:** In a **push** model (like RabbitMQ), the broker actively sends messages to consumers. In a **pull** model (like Kafka), consumers request messages from the broker when they are ready.
 - **Fan-out:** A pattern where a single message is delivered to multiple consumers. This is useful for tasks like cache invalidation, where multiple cache nodes need to be updated simultaneously.
 - **Message Persistence:** The queue's ability to store messages on disk. This ensures that messages are not lost if a consumer or the broker itself restarts.

Service & Deployment Management

Service Management involves maintaining the health and lifecycle of your applications.

- **Version Control:** Using tools like Git to track changes to your codebase.
- **JAR Handling:** Managing Java Archive (JAR) files, which package your application. This includes versioning them and storing them in an artifact repository (like Nexus or Artifactory).
- **Backward Compatibility:** Ensuring that new versions of a service can still work with older clients or services. This is crucial for avoiding breaking changes in a distributed system.

Deployment Strategies are methods for releasing new code to production.

- **Canary Deployment:** Rolling out the new version to a small subset of users first. If no errors occur, you gradually roll it out to everyone.
- **Blue-Green Deployment:** Setting up a new, identical production environment ("Green") with the new version. Once it's tested, you switch traffic from the old environment ("Blue") to the Green one. This allows for instant rollback if issues arise.

- **Manual vs. Automated Scaling:** **Manual scaling** is when an operator adds or removes servers based on observation. **Automated scaling** (or auto-scaling) is when the system automatically adjusts the number of servers based on predefined metrics like CPU usage or request count.

Deployment Management refers to the practicalities of deploying in a cloud environment.

- **GCP/AWS Integrations:** Using cloud provider tools for CI/CD, like AWS CodePipeline or Google Cloud Build, to automate the build, test, and deployment process.
- **Project Structuring:** Organizing your cloud resources (e.g., using separate GCP projects or AWS accounts for development, staging, and production) for better security, billing, and management.

Containerization & Infrastructure

Docker & Kubernetes are the foundation of modern cloud-native applications.

- **Docker:** A tool to package an application and its dependencies into a portable unit called a **container**.
- **Kubernetes:** A **container orchestration** platform that automates the deployment, scaling, and management of containerized applications across a cluster of machines. It handles networking between containers and can manage setups like a Spring Boot application connecting to a MySQL database, both running in separate containers.

Internet Infrastructure includes the core protocols that make the internet work.

- **DNS (Domain Name System):** The internet's phonebook, which translates human-readable domain names (e.g., `google.co.in`) into machine-readable IP addresses (e.g., `142.250.196.195`).
- **Anycast:** A networking technique where the same IP address is assigned to multiple servers in different geographic locations. When a user sends a request, it's automatically routed to the nearest server, reducing latency.
- **BGP (Border Gateway Protocol):** The protocol that manages how packets are routed across the internet. It helps different networks (like your ISP and Google's network) find the best path to send data to each other.

CDNs (Content Delivery Networks) are networks of globally distributed **edge servers**. They cache static content (like images, videos, and CSS files) closer to users, which dramatically reduces page load times.

System Robustness & Common Challenges

System Robustness is about designing systems that can withstand failures.

- **Handling Dependencies:** A service should be designed to handle failures in its dependencies gracefully. For example, if a payment service is down, an e-commerce site should disable the checkout button rather than crashing the entire site.
- **Failure Points:** Identifying and mitigating single points of failure (SPOFs) in your architecture. This often involves adding redundancy, like running multiple instances of a service or database.

Common Challenges in distributed Java environments include:

- **JVM DNS Caching:** The Java Virtual Machine (JVM) caches DNS lookups. If the IP address of a service changes, the JVM might not pick up the change without a restart, leading to connection failures.
- **Proxy Configs:** Ensuring that applications running inside a corporate or cloud network are correctly configured to route outbound traffic through a proxy server if required.
- **Port Mapping:** In containerized environments (like Docker), you must map a port from the host machine to a port inside the container to allow external traffic to reach your application.